

NAMEI_EXT: Programmable VFS Name Resolution without Filesystem Ownership

Anonymous

Abstract

Modern agent, build, and service systems increasingly construct task-specific filesystem views around ordinary tools. Binding a pathname to the wrong lower object can corrupt an agent workspace, reuse stale build state, or expose an unintended service configuration. The root cause is a boundary mismatch: existing systems couple runtime pathname binding either to namespace materialization before lookup or to FUSE, stackable, or custom filesystem ownership, even when ordinary file semantics should remain with the lower filesystem. We argue that dynamic filesystem views are a late-binding problem, not necessarily a new-filesystem problem: runtime policy should select an existing VFS object, while the VFS and lower filesystem retain object ownership and native filesystem semantics. The challenge is that an in-progress path walk is not an object pointer: it carries operation intent, root and mount constraints, RCU/ref-walk state, and sequence validation. The kernel therefore needs a narrow contract for accepting a policy-selected object without bypassing those constraints or transferring filesystem ownership. NAMEI_EXT places one bounded eBPF decision function at name resolution, represents selectable objects by opaque IDs for kernel-held paths, resumes existing VFS processing after validation, and bypasses unrelated paths with an early scope filter. Six of the seven required source-controlled or source-derived workflows, including official `xdg-document-portal` 1.18.4, pass reviewed KVM oracles. In the headline Agent case, two concurrent process groups select different Click workspace roots at one pathname; a released task observes 39/40 tests on the base view and 40/40 on the completed view, while switch and rollback change object identity and task outcome. Against feature-equivalent FUSE, the Agent lifecycle takes $5.51\mu\text{s}$ versus $62.64\mu\text{s}$, and cache-hot FxMark `SELECT` reaches $1.052\text{--}1.088\times$ FUSE throughput; a separate 60-boot experiment meets its patched-unattached fast-path criterion. Directory enumeration reaches $2.20\text{--}3.66\times$ FUSE throughput in five of six cells, and a matched Wraps-derived experiment passes 37/37 workspace oracles and 21/21 fault cells in each of three boots while exposing the broader filesystem-method boundary.

1 Introduction

Modern systems increasingly construct task-specific filesystem views around ordinary tools. Agents branch, fork, checkpoint, hide side effects, and run build or test commands over repository workspaces [46, 26, 51]. Environment systems validate repositories against generated Docker state, cache epochs, and verified local artifacts [12, 11, 13, 34], while services consume projected configuration, secrets, certificates, or fixtures as operational state changes [19, 20, 18].

However, these views make pathname binding a correctness boundary. Binding a name to the wrong lower object can validate the wrong revision, leak a write into a parent workspace, feed stale cache input into a build, or expose the wrong service configuration. In many of these cases, the needed behavior is not a new file format, storage service, or custom file operation. The needed behavior chooses which existing object a name denotes, or whether that object is visible.

The root cause is that existing mechanisms couple pathname binding to a broader owner. Namespace construction binds names before use by materializing or updating mounts, layers, links, or files. FUSE, stackable, and custom filesystems can instead choose a binding during lookup, but that choice lives inside a filesystem implementation that also owns a wider operation and lifetime boundary. Workload state may need to change only which existing object a name denotes, while data, writes, permissions, page-cache behavior, persistence, and consistency should remain with the lower filesystem.

Existing mechanisms bracket this lookup-time selection problem. Namespace mechanisms such as bind mounts, OverlayFS, projected volumes, and copied trees construct or materialize a namespace before the application uses it [23, 43, 19]. eBPF LSM hooks can attach verified access-control policy to security decisions [42], but they are not an object-selection boundary. FUSE, stackable filesystems, custom filesystems, and metadata services can compute richer views, but they cross into a broader filesystem boundary [38, 52, 25, 54, 31, 30]. They may be the right mechanism when the oracle needs synthetic contents, custom metadata, data-path mediation, or runtime orchestration, but they are broader than lookup-time name-resolution policy.

We argue that dynamic filesystem views are a pathname late-binding problem, not necessarily a new-filesystem problem. The recurring subproblem is a state-dependent path view: workload state changes which existing lower object a pathname denotes or whether that object is visible. VFS name resolution is where this binding is chosen. The key separation is therefore to make binding policy programmable while the VFS and lower filesystem retain object ownership and ordinary filesystem semantics.

The central challenge is defining when the VFS can accept a policy-selected object as a normal name-resolution result. An in-progress path walk is not an inode pointer: `nameidata` carries operation intent, the current root and mount, traversal restrictions, RCU/ref-walk state, and sequence numbers used to validate concurrent rename and mount activity. Replacing only its path can bypass those constraints or leave the walk with inconsistent ownership. The extension must therefore admit only pre-authorized, type- and operation-compatible objects; obtain stable lifetime without exposing kernel pointers to BPF; reject selections whose path constraints cannot be preserved; and then resume the original VFS state machine rather than emulate it.

Two additional requirements follow from where this object-selection boundary sits. Intermediate components, final open/create components, and directory entries traverse different VFS paths, but must implement the same binding contract. Cache-hot lookup performs that work during lockless RCU-walk, while target registration and replacement occur concurrently. Finally, the decision point lies on a system-wide hot path, so disabled and unrelated lookups must bypass policy work before context construction.

We present `NAMEI_EXT`, which separates pathname-binding policy from filesystem ownership. One eBPF function returns bounded pass, redirect, hide, or `SELECT_TARGET` actions at the VFS paths that observe names. Opaque target IDs refer to kernel-held `struct path` objects; the kernel validates each result, converts an RCU-borrowed target into stable `namei` state, and resumes ordinary permission, open, exec, and lower-filesystem operations on the selected object. A static key and an RCU parent-scope filter bypass policy work before context construction when no relevant policy exists. The resulting ownership split is analogous to `sched_ext`: policy is programmable, while the kernel retains the subsystem machinery that owns objects and executes semantics [37].

We evaluate whether this ownership split is sufficient for source-derived path-view policies, cheaper than feature-equivalent FUSE placement, and narrower in method and runtime responsibility than broader filesystem ownership. RQ1 asks whether KVM workloads pass the same correctness conditions while preserving lower-filesystem semantics. Reviewed formal KVM runs pass fixed oracles for Agent workspaces, official XDG Documents portal application-sharing grants, concurrent Bazel action views, the already-materialized payload-view subset of Kubernetes `AtomicWriter`, and toolchain environment selection. The headline Agent result includes three fresh boots of a released Click source task: concurrent base and completed views produce 39/40 and 40/40 tests, and switch/rollback changes both the selected object and the task outcome. RQ2 compares the same policies with feature-equivalent FUSE: the matched Agent lifecycle is $11.32\times$ [11.24, 11.64], and cache-hot FxMark `SELECT` reaches $1.052\text{--}1.088\times$ cached-FUSE throughput across MRPL, MRPM, and MRPH. On corrected FxMark directory enumeration, `SELECT` reaches $2.20\text{--}3.66\times$ FUSE throughput in five of six cells; at four shared-directory workers, the ratio is 1.018 [0.907, 1.135], exposing the shared-directory contention boundary. RQ3 executes a matched Wraps-derived implementation; both mechanisms pass 37/37 workspace oracles and all 21 fault cells in each of three boots, while runtime tracing exposes the broader stackable filesystem-method boundary.

The paper’s primary contribution is the design and Linux implementation of a late-binding boundary: programmable selection of existing objects during VFS name resolution while preserving VFS and lower-filesystem ownership. The paper makes two concrete contributions:

- the design and modified-Linux implementation of a pathname-binding extension, including one contract across component lookup, final lookup, and `readdir`, registered-target lifetime across RCU/ref-walk,

preservation of normal VFS processing after selection, and a scope-aware fast path (Sections 4 and 5); and

- an evaluation that holds workload correctness fixed while comparing NAMEI_EXT with feature-equivalent FUSE and recording which responsibilities would move into custom or stackable filesystems (Section 6).

2 Background

2.1 VFS Name Resolution

The Linux VFS resolves a pathname by walking directory components, consulting directory entries and inodes, and then dispatching ordinary file operations to the selected filesystem [44, 45]. Lookup selects the object on which later file semantics operate, while the VFS and lower filesystem retain permission checks, reads and writes, page-cache behavior, persistence, and consistency. The lookup/operation separation is the boundary the paper uses for path-view policy.

2.2 Programmable Kernel Policy

eBPF lets the kernel run verified, bounded programs at explicit attachment points [35, 36]. The `sched_ext` scheduler class is a useful precedent because it lets BPF choose scheduling policy while the kernel retains scheduler machinery [37]. NAMEI_EXT applies the same policy-versus-ownership separation to VFS name resolution: BPF supplies a bounded path-view decision, and the kernel retains the VFS code that owns dentries, inodes, and object lifetime. The analogy concerns the ownership split rather than the scheduler interface.

2.3 Namespace Construction And Filesystem Services

Linux already offers mechanisms for alternate filesystem views. Bind mounts, OverlayFS, projected volumes, and copied trees construct or materialize a namespace before the application uses it [23, 43, 19]. FUSE and stackable or custom filesystems provide a broader programmable filesystem boundary, often by placing a service or filesystem implementation on the operation path [38, 52, 25]. These mechanisms define the neighboring boundaries for the path-view policy studied in the rest of the paper.

3 Motivation And Workloads

A VFS name-resolution extension targets a narrow workload pattern in which the source system changes which existing object a pathname should select, but the operation does not need a new filesystem object model. The source-system evidence below identifies that pattern and separates it from the broader behavior that should remain with agent runtimes, environment builders, services, FUSE filesystems, or custom filesystems.

3.1 State-Dependent Path Views

Under this definition, a transition qualifies when a system state change alters which existing object a path-name denotes, or whether the object is visible, while ordinary lower-filesystem semantics remain unchanged. The definition is intentionally narrower than "programmable filesystem." It excludes synthetic file contents, custom metadata persistence, distributed directory indexing, data-path mediation, write-conflict resolution, and cross-path transactional snapshots.

Within this definition, the transition is the workload state change, the policy is the eBPF decision logic, and the action is the bounded result that the kernel validates for lookup or directory enumeration.

The distinction matters because the natural implementation boundary changes. If the system needs synthetic contents or custom metadata persistence, a FUSE service, custom filesystem, metadata service, or application/runtime mechanism is the appropriate implementation boundary. If the state change only affects

Example	State and operation	NAMEI_EXT action and delegated behavior	Oracle
Agent workspace lifecycle	branch or checkpoint state changes lookup and readdir under a workspace path	select or hide existing lower objects while the source runtime keeps COW, checkpoint, sandbox, writes, and lifecycle orchestration	final tree, lookup/readdir coherence, and preserved lower-filesystem permissions
Application file sharing	application identity and grant state change whether one existing host document is visible	select the granted document or hide it while grant persistence and portal UI remain outside the hook	grant, revoke, cross-application isolation, and unchanged lower object
Build action sandboxing	action identity and declared-input set choose which existing input tree one logical pathname exposes	select the action’s declared object or hide an undeclared name while Bazel keeps process and build semantics	real action success, distinct concurrent outputs, and undeclared-input failure

Table 1: Representative state-dependent path-view examples. The NAMEI_EXT evidence is the lookup/readdir action, and non-name-resolution behavior remains with the source runtime or lower filesystem.

lookup or directory visibility, a VFS name-resolution policy can express the relevant policy. The remaining source-system behavior stays outside name resolution.

3.2 Workload Sources

We use source systems to choose workloads and correctness conditions, not as separate contributions. We call each extracted correctness condition a source oracle. A same-oracle KVM run applies that condition unchanged to NAMEI_EXT and its comparison rows.

The first source family is agent/workspace state. AgentFS is the primary source because its public implementation exposes workspace lifecycle traces from which path-selection and visibility oracles can be extracted [46]. SWE-Factory-Gym supplies a released Click source task whose fail-to-pass test can run through those workspace views [11, 33]. BranchFS, Sandlock, and YoloFS provide neighboring branch, sandbox, staging, and stackable-filesystem context rather than separate evaluation workloads [26, 27, 51].

The traditional source families are application file sharing, build action sandboxing, checkpoint/restore, HPC staging, service rotation, and toolchain environments. The application-sharing workload follows the XDG Documents portal grant/revoke behavior [48]. The build-action workload follows Bazel’s action-specific input views [3]. The toolchain workload follows per-environment software selection and switch/rollback behavior represented by profile systems such as Nix [28]. Kubernetes AtomicWriter supplies the completed already-materialized payload-view sequence [19]. DMTCPath path virtualization supplies an additional source-defined lifecycle shape but remains a portfolio case until its exact KVM oracle runs [2]. Full projected-volume materialization and service reload also remain outside the completed Kubernetes subset.

A historical ccache matrix supplies supporting compile-time evidence, but is not a headline source workload: ccache already implements cache lookup and validation in userspace. Source systems are admitted only when the behavior under study is the application’s path view, not an application-level cache algorithm re-expressed in BPF.

Only source behavior tied to a concrete same-oracle KVM run becomes final evaluation evidence. Source inspection and paper-derived behavior motivate workload shape and boundary comparisons.

Three design constraints follow from the source systems. First, each workload carries its own state, such as a branch identifier, checkpoint, cache epoch, build generation, service epoch, or secret epoch. Policy therefore must run at the affected lookup or directory-enumeration operation and remain scoped per workload.

Second, the path-view portion of each source system is narrower than the full system. Many rows also include runtime orchestration, service reload, synthetic contents, or storage semantics that remain separate responsibilities. The narrower path-view scope implies that the kernel and lower filesystem must keep data, permission, and persistence ownership.

Third, directory visibility and lookup selection appear together in the source oracles. Lookup and readdir therefore need one coherent decision contract.

Together, these source-system observations set the paper’s comparison rule: (1) source behavior supplies the oracle, (2) feature-equivalent FUSE supplies the direct programmable-policy cost comparison, and (3) custom or stackable filesystems define the broader filesystem-method and runtime-responsibility boundary.

Materialized namespace mechanisms remain related-work and background comparisons rather than the central experiment. The same rule drives the design requirements in Section 4 and the RQ-organized evaluation in Section 6.

3.3 Selected Workloads

The RQ1 evaluation set comprises seven non-overlapping industrial workflows: W1 Application File Sharing, W2 Agent Workspaces, W3 Build Action Sandboxing, W4 Service Configuration and Secret Rotation, W5 Checkpoint/Restore and Migration, W6 HPC File Staging, and W7 Toolchain and Dependency Environments. All seven are required RQ1 cases; RQ2 and RQ3 use representative cases for deeper cost and ownership comparisons.

W1 replays the per-application grant, isolation, and revoke behavior of `xdg-document-portal`. W2 exercises branch-visible staged/base selection, hides, `readdir` coherence, and lower-tree mutations, then runs a released Click task through concurrent base/completed views, switch, rollback, and withdrawal. W3 executes concurrent Bazel actions with distinct declared-input roots and an undeclared-input failure. W4 follows the Kubernetes `AtomicWriter` payload-view sequence through update, no-op, and rollback under one stable logical root. W5 uses DMTCP’s moved-root restart and reopen behavior. W6 uses Spindle’s shared-versus-node-local object selection during HPC program launch. W7 selects two existing CPython environments, switches one process-group view, and rolls it back.

All seven workflows have reviewed KVM source-oracle results. A workload becomes evidence only after its reviewed KVM source-oracle run. In every completed case, policy runs at lookup and directory enumeration; non-name-resolution behavior stays with the source system or lower filesystem.

4 Design

4.1 System Overview

`NAMEI_EXT` treats a dynamic filesystem view as late binding of a pathname or directory entry to an existing VFS object. Existing systems commonly couple that binding either to namespace materialization before lookup or to a filesystem implementation that owns lookup and subsequent operations. `NAMEI_EXT` separates the binding policy from filesystem ownership: policy chooses a bounded binding, while the VFS validates the result, owns the selected object, and executes ordinary lower-filesystem semantics.

Three invariants organize the design. First, every VFS path that observes a name implements the same bounded binding contract. Second, policy returns only names or opaque handles; the kernel alone validates object type, scope, and lifetime before installing a result. Third, a selected binding changes the object on which ordinary VFS processing operates, not the semantics of those operations. Permissions, file operations, data, writes, page cache, persistence, and consistency remain with the VFS and lower filesystem. A separate deployment requirement keeps this optional boundary off disabled and unmanaged paths before policy-context construction.

The central design question is the acceptance contract for a policy-selected object. Native name resolution does not return a bare inode: the surrounding `nameidata` records how the object was reached, which operation is being performed, which root and mount restrictions apply, whether the walk is under RCU, and which sequence values protect it from concurrent namespace changes. `NAMEI_EXT` may replace the visible object only when it can preserve the relevant state or fail the operation before the result is used. This contract, rather than the BPF invocation itself, determines the mechanism boundary.

Figure 1 shows the resulting flow. An application in a workload scope performs an ordinary lookup, open, stat, exec, or `readdir`. The VFS reaches the `NAMEI_EXT` decision point before committing to the visible object for the current component or directory entry. The eBPF policy returns a bounded path-view action, and ordinary VFS and lower-filesystem logic resumes after kernel validation.

4.2 One Binding Contract Across Resolution Paths

Linux does not resolve every name through one function. Intermediate components pass through the component walk, the trailing component of `open()` has separate create and open handling, and `readdir()` obtains




Figure 1: NAMEI_EXT separates late-binding policy from filesystem ownership. BPF chooses a bounded binding; the kernel validates and installs it according to the surrounding component, final-open, readdir, and RCU/ref-walk path; ordinary VFS semantics then operate on the selected object.

names from a filesystem iterator. NAMEI_EXT applies one policy contract at all three points rather than assigning each call site an independent rule language. The input identifies a lookup or directory-entry event; policy state and the current parent/name pair determine the result.

The action meanings are consistent across those paths. PASS exposes the ordinary lower result. HIDE returns absence from lookup and suppresses the corresponding directory entry. REDIRECT substitutes a validated component name for lookup or for the emitted directory entry. SELECT_TARGET selects a registered existing object during lookup; directory iteration over a selected directory proceeds after that directory has been opened. The current design does not use SELECT_TARGET to invent directory entries, because synthetic aliases require an object model and stable iterator positions beyond existing-entry filtering.

The ABI gives lookup and readdir one action vocabulary, but the kernel invokes the two event types independently. A policy is responsible for making its lookup result and directory view coherent; workload oracles test that obligation. Target selection changes the object on which subsequent ordinary VFS processing operates. Create semantics remain outside selection: REDIRECT and SELECT_TARGET reject LOOKUP_CREATE, and the policy does not decide target creation or copy-up behavior.

The contract linearizes each decision at one program invocation. It does not provide a transaction across independent system calls or across all entries of a directory stream. A controller can publish immutable target generations and switch one policy-state selector atomically, as the evaluated workloads do, but policies that require cross-name transactional updates need a broader mechanism. Stating this update model separates binding coherence for a fixed policy state from filesystem transaction semantics that NAMEI_EXT does not own.

4.3 Accepting A Policy-Selected Object

A decision is not yet a usable VFS object. BPF cannot return a kernel pointer or an arbitrary path string, and cache-hot RCU-walk does not hold ordinary references that a policy can transfer. NAMEI_EXT therefore separates policy output from object ownership. A control plane registers an existing object, the kernel retains its `struct path`, and policy selects it by an opaque ID. The kernel validates the action, target type, scope, and operation before installing the result in the current namei state.

Object handoff differs between ref-walk and RCU-walk without changing policy meaning. Ref-walk acquires an independent path reference directly. RCU-walk borrows the registry’s kernel-held path, then uses ordinary namei sequence validation to obtain stable references or restart if concurrent replacement invalidates the walk. Policy chooses only the binding; object lifetime and the transition from a speculative walk to stable VFS state remain kernel responsibilities.

The accepted subset is defined by five checks. First, registration from an open descriptor establishes both controller authority and kernel-held target lifetime; policy may name only the resulting cgroup-scoped ID. Second, a target must be type-compatible with its position: an intermediate target is a directory, while a final target is an existing regular file or directory. Third, selection cannot own creation, so LOOKUP_CREATE is rejected. Fourth, selection fails when an `openat2()` scoped walk could lose its root constraint, or when

Goal	Design choice	RQ
G1 Expressiveness	invoke one workload-scoped decision contract at lookup and directory-enumeration time	RQ1
G2 Semantic preservation	select only existing lower objects and resume ordinary VFS processing for data, writes, permissions, page cache, and persistence	RQ1, RQ3
G3 Placement cost	run bounded policy at lookup/readdir instead of routing the same decision through a filesystem service	RQ2
G4 Output validation	use verified eBPF outputs plus kernel validation so malformed or unsupported bindings fail visibly	Implementation correctness

Table 2: Design-goal mapping. If an oracle needs synthetic contents, custom metadata, or data-path mediation, the effect belongs outside NAMEI_EXT’s boundary.

RESOLVE_NO_XDEV would cross a mount. Fifth, RCU-walk may borrow a registered path only until ordinary namei validation can stabilize it; otherwise the walk restarts or fails. Unsupported cases do not fall back to the unmodified lower name.

Registration is therefore an authorization boundary, not merely a lookup cache. The controller authorizes an already opened object for selection by one policy scope. Subsequent VFS processing checks the logical walk and selected object, but does not replay traversal through the target’s physical parents. Policies that require unprivileged arbitrary-path selection or scoped traversal through a replacement path need a broader mechanism.

4.4 Preserving Filesystem Semantics

After installing a selected object, NAMEI_EXT resumes the surrounding VFS state machine rather than starting a second path walk or forwarding filesystem methods. Intermediate directory selection continues component walking from the selected directory; final selection continues permission, stat, access, open, or exec completion on the selected object; opening a selected directory uses its lower iterator. An already opened file or directory subsequently executes lower-filesystem operations without reinvoking policy.

This preservation invariant is what keeps NAMEI_EXT from becoming a filesystem. The extension does not emulate permission checks, mount handling, dentries, inodes, reads, writes, page cache, persistence, or filesystem-specific behavior. For any operation after a successful binding, the required result is the result of ordinary VFS execution on the selected lower object.

4.5 Name-Resolution Boundary Goals

Four evaluation goals follow from this ownership split: expressiveness, semantic preservation, placement cost, and visible output validation. Like `sched_ext`, NAMEI_EXT lets BPF choose policy while the kernel keeps the VFS mechanisms that own objects and execute filesystem semantics.

The design defines a narrow kernel extension for name-resolution policy while leaving full filesystem expressiveness to FUSE, stackable, or custom filesystems. Table 2 maps the four design goals to the RQs.

4.6 Fast-Path Contract

The extension sits on a system-wide hot path, so its cost model is part of the interface rather than an implementation afterthought. When no program is attached, lookup and directory iteration must follow the ordinary VFS path after one static-key test. When a program is attached to a limited set of parents, an unrelated parent must be rejected before policy-context construction or workload-scope discovery. This first test is conservative: false positives may execute the authoritative scope check, but a false negative would bypass policy and is forbidden.

For managed parents, cache-hot lookup may already be in RCU-walk. The decision contract therefore cannot require a sleeping lock or an owned dentry/mount reference. NAMEI_EXT separates the BPF decision from action-specific object handling: pass-through can remain in RCU-walk, while an action that needs refcounted state must use normal namei conversion and validation. Walk mode is not exposed as workload policy, so the same pathname and policy state do not produce different decisions merely because the VFS cache changed.

A successful in-place conversion applies the saved decision without calling the program again. If conversion fails and `ordinary namei` returns `ECHILD`, however, the VFS may restart the complete pathname walk and execute policy again. The interface therefore provides at-least-once execution across full VFS restarts, not exactly-once execution. Policies must keep externally visible side effects idempotent; map reads and writes that define the decision remain subject to the same restart rule.

4.7 Policy Model

The workloads require policy to see lookup and directory-enumeration events, carry workload state, return only bounded path-view actions, and let the kernel validate the result.

The policy input contains operation context, the current component, parent context, and policy state. The policy state is workload-defined but not a filesystem object model. For the representative workloads, policy state is a branch or checkpoint identifier in an agent workspace, a cache or validation epoch in an environment run, or a configuration/secret epoch in a service.

The output passes the lower lookup, redirects selection to a replacement component, hides a name from lookup and directory enumeration, or selects a registered lower object. Access-control denial remains outside the model because `NAMEI_EXT` focuses on selecting or hiding existing objects during name resolution. Directory enumeration uses the same decision function and action vocabulary as lookup; the policy must implement the corresponding direction, and evaluation checks the resulting lookup/readdir view.

The policy cannot return a kernel pointer or an arbitrary pathname. A control plane first registers an open existing object, and the kernel retains its `struct path` under an opaque ID. The policy may return only that ID. Intermediate selections must identify directories; final selections may identify existing regular files or directories. The selected object then continues through normal `namei` validation, permission checks, open completion, and lower-filesystem file operations.

The kernel also validates action type, name bounds, target identity, selected object type, create restrictions, and path-walk scope. Invalid decisions fail visibly rather than falling back to a weaker policy. This registered-object model bounds what policy can select without turning the hook into arbitrary recursive path rewriting.

4.8 Filesystem Ownership Boundary

`NAMEI_EXT` excludes filesystem ownership by leaving inode operations, file operations, dentry and inode allocation, synthetic contents, custom metadata, read/write mediation, distributed indexes, and cross-path transactions outside the hook.

The name-resolution boundary helps only when the source transition can be expressed as lookup selection or directory visibility. When the oracle requires data-path mediation, synthetic content, custom metadata, or application-level orchestration, the transition remains motivation or related work rather than `NAMEI_EXT` evidence.

5 Implementation

The Linux prototype turns the policy contract into a real `cgroup/namei_ext` attachment path. Policies are eBPF programs under `bpf/policies/*.bpf.c`. The implementation consists of the VFS call sites, the BPF ABI, user-space policy attachment and map setup, and kernel-side validation for the supported action set.

5.1 Path-Walk Integration

The prototype integrates policy at three VFS points. `walk_component()` handles intermediate and ordinary trailing components, `open_last_lookups()` handles the final component for open/create, and `iterate_dir()` mediates entries returned by a lower directory iterator. Each call site first tests the global enable key and parent-scope filter. A matching lookup constructs the common context and invokes the same `cgroup/namei_ext` decision function before the VFS commits to a lower result.

The component and final-open wrappers preserve the surrounding `namei` state rather than starting a second path walk. `PASS` enters the ordinary component or final-open implementation. `REDIRECT` replaces

only the current `qstr`, invokes the lower directory's `d_hash()` operation, and restores the original component after lookup. `HIDE` returns `ENOENT`. `SELECT_TARGET` installs a validated registered path into the existing `nameidata`; intermediate walking or final VFS completion then continues from that object.

Cache-hot lookup reaches these wrappers in RCU-walk. The BPF-visible flags mask the internal `LOOKUP_RCU` bit so cache state does not change the policy input. `PASS` remains in RCU-walk. Same-parent redirect converts through `try_to_unlazy()` before changing the component, while errors preserve the VFS conversion and `RESOLVE_CACHED` behavior. The wrapper does not reinvoke the program merely to apply a saved action after successful conversion. If conversion returns `ECHILD`, the surrounding VFS may restart the complete pathname walk and invoke policy again. Programs therefore cannot assume exactly-once execution across VFS restart and must make observable side effects idempotent.

5.2 BPF ABI

The kernel-facing ABI exposes one decision function. The context includes an event type, operation flags, the current component name, parent identity, identifiers used by the policy, and bounded output fields for the selected path-view action. The current action set is `PASS`, `REDIRECT`, `HIDE`, and `SELECT_TARGET`. `REDIRECT` carries a bounded replacement component in `redirect_name`. `SELECT_TARGET` carries an opaque `target_id` that indexes a kernel-registered lower `struct path`. The policy cannot return arbitrary unbounded path strings or perform recursive path walks.

The verifier exposes names and parent identity as read-only context and permits writes only to the bounded replacement-name and target-ID outputs. After program execution, the kernel validates action type, name length and contents, target identity, object type, and selection scope before continuing. Redirect decisions reject empty names, `.`, `..`, embedded nulls, slashes, and create operations.

5.3 Registered-Target Lifetime

Workload controllers open each target with `O_PATH` and register its file descriptor under an opaque target ID. The kernel copies the descriptor's `struct path`, holds mount and dentry references, and publishes the record in an RCU hash table keyed by the policy scope and ID. Replacing or clearing a record uses two related paths: replacement atomically publishes the new RCU hash entry, while clearing removes the old entry. Both wait for an RCU grace period before dropping the retired path references. A policy never receives the path or a kernel pointer.

Ref-walk target selection looks up the record and acquires independent path references. RCU-walk instead borrows the registry's references while the outer `namei` RCU critical section protects the record. When `namei` installs a borrowed target, it samples the target dentry sequence and retains the mount sequence. Existing `try_to_unlazy()` and `complete_walk()` validation either obtain stable references or reject and restart a changed walk. Thus concurrent target replacement cannot free a path still borrowed by an RCU reader.

Before installation, the kernel rejects missing IDs, symbolic links, unsupported object types, create-through-select, scoped walks, and prohibited cross-mount selection. Intermediate targets must be directories. A final target may be an existing regular file or directory, after which ordinary `stat`, `access`, `open`, `exec`, `O_PATH`, or `O_DIRECTORY` handling proceeds on the selected object. The prototype does not create targets, copy data, populate caches, or implement lifecycle orchestration.

5.4 Replacement-Component Validation And Permission Preservation

Workload setup configures the policy state that chooses replacement components. The prototype implements the redirect subset by copying a bounded `redirect_name` from the BPF output, validating its length, action, and same-parent scope, and then resuming the normal VFS path. The replacement component is not an arbitrary path string and does not trigger a recursive path walk. After validation, the normal VFS path continues, including permission checks and lower-filesystem file operations. Same-parent validation keeps `NAMEI_EXT` from becoming arbitrary path rewriting or a filesystem service. Both component replacement and registered-target selection reject `LOOKUP_CREATE`: they may choose an existing object, but they cannot change lower-filesystem creation or copy-up rules.

The implemented hide action uses the same decision point but does not allocate or synthesize filesystem objects. Lookup observes absence, and directory enumeration suppresses the lower entry for the active policy state.

5.5 Directory Iteration

`iterate_dir()` wraps the lower filesystem’s existing `dir_context.actor`. For every returned lower entry, the wrapper invokes the common decision function. `PASS` forwards the original name and inode information, `REDIRECT` forwards a validated replacement name, and `HIDE` consumes the lower entry without exposing it. The wrapper copies the iterator position, count, and directory-entry flags in both directions so partial `getdents64()` calls retain the lower iterator’s progress.

`SELECT_TARGET` is intentionally rejected as a per-entry `readdir` action. Selecting an existing directory during lookup and then enumerating that opened directory is supported; inventing a new alias for an arbitrary target is not. The latter would require synthetic-entry storage and stable position semantics that the current existing-object boundary does not own.

5.6 Unused And Unmanaged Fast Paths

The cgroup BPF attachment uses a static key for the globally disabled case. With no active program, component lookup, final-open lookup, and directory iteration enter their ordinary implementations after that test. An attached program may additionally register up to eight exact parent paths. An RCU hash index and active-global counter let all three call sites reject an unrelated parent before entering a noinline wrapper, constructing a BPF context, or discovering the current cgroup.

Scope updates run under the scope mutex, while a raw spinlock sequence counter publishes the global counters, RCU hash nodes, and active state without a sleeping lock in RCU-walk. Removed scope records and held path references are released only after an RCU grace period. The prefilter is conservative: a positive result still executes the effective-owner and per-policy scope checks, while counter overflow saturates toward policy execution rather than creating a false negative.

5.7 Failure Handling

Verifier-invalid programs cannot attach. At runtime, an unknown action, malformed redirect, missing registration, incompatible target type, or unsupported operation returns a specific error to the pathname operation. The implementation does not translate those errors into `PASS`; doing so would expose the ordinary lower name when the requested policy did not execute. This visible error behavior and the unchanged lower object after invalid decisions are implementation-correctness checks; RQ3 primarily compares which filesystem methods and runtime state each mechanism owns.

Phase 1 validation runs in KVM with the modified kernel and the real `cgroup/namei_ext` attach path. Host-only checks, object inspection, and `bpftool`-only smoke tests are diagnostics rather than validation for the paper’s mechanism claims.

These components give the evaluation a single validation path in which Make targets boot the modified kernel, attach policies through `cgroup/namei_ext`, and preserve raw validation and experiment artifacts.

6 Evaluation

6.1 Research Questions

The evaluation is organized around three research questions. Each answer is conditioned on a source oracle: the same workload behavior must be used for `NAMEI_EXT`, the feature-equivalent FUSE row, and the custom/stackable responsibility comparison. One-off mechanism checks, source surveys, and weak baselines do not answer an RQ by themselves.

RQ1 Can a narrow VFS name-resolution extension express real state-dependent path-view policies without taking over filesystem semantics?

RQ2 What is the cost of putting programmable policy on the VFS name-resolution path compared with a feature-equivalent FUSE policy implementation?

RQ3 When the required policy is only name resolution, which filesystem methods and runtime responsibilities does NAMEI_EXT own compared with a custom or stackable filesystem?

6.2 Experimental Setup

All NAMEI_EXT mechanism runs execute in KVM [40] with the modified kernel and the real `cgroup/namei_ext` attach path. Make targets own the run, attach the policy, collect raw artifacts, and fail on missing capability, verifier failure, syscall failure, or oracle failure. Each final row records the kernel commit, image identity, policy object, workload inputs, stdout/stderr, dmesg, per-operation lookup/readdir events, and comparison-specific setup observations. Operation-weighted metrics weight each policy action by its observed frequency in the workload trace.

Each final reported run records the kernel configuration and image identity, filesystem configuration, run count, warmup rule, confidence-interval method, cache protocol, compared implementation, source commit, and raw observations.

Correctness gates every number. A FUSE row is used for RQ2 only after it implements the same policy over the same source oracle. No-hook and lower-filesystem rows calibrate fixed hook and measurement cost. They are not competing baselines. RQ3 executes a matched Wraps-derived implementation for the Agent workspace oracle and compares filesystem-method and runtime responsibility rather than throughput.

6.3 RQ1: Expressiveness And Sufficiency

RQ1 asks whether the VFS boundary can express path-view transitions extracted from source oracles while preserving lower-filesystem semantics. A reportable row must pass the source oracle through the real KVM attach path and show that NAMEI_EXT changed only lookup or directory visibility, while runtime orchestration, storage state, synthetic contents, writes, and data-path behavior remain with the source system or lower filesystem.

Workload	Source oracle	NAMEI_EXT evidence	Result evidence
Application file sharing (supporting) [48]	Official <code>xdg-document-portal</code> 1.18.4 executes <code>Add</code> , <code>grant</code> , cross-application isolation, and <code>revoke</code> through its FUSE view.	NAMEI_EXT selects or hides the same existing-object view by application and grant state.	Three fresh KVM boots passed 15/15 official-source and 15/15 NAMEI_EXT states with exact operation and enumeration agreement. The NAMEI_EXT visible path matched its registered lower object; both mechanisms preserved their direct lower object.
Agent workspaces [46, 11, 33]	AgentFS-derived branch state selects staged/base objects and hides whiteouts; a released Click task distinguishes base from completed source.	NAMEI_EXT selects or hides existing lower objects; the lower filesystem retains permissions, writes, metadata, and file data.	The lifecycle matrix recorded 1,170 pass-bearing records with zero failed oracles plus six metadata records in each of three runs. Three fresh source-task boots passed 12/12 policy-backed task states and 6/6 physical source controls: concurrent base/completed views produced 39/40 and 40/40 tests, followed by switch, rollback, and withdrawal. Three fresh KVM boots completed all six Bazel 6.5.0 actions, matched six logical/lower input objects, hid undeclared inputs, and preserved 12 lower objects.
Build action sandboxing (traditional main case) [3]	Two concurrent Bazel actions use one logical input path but different declared roots; an undeclared existing input must be absent.	NAMEI_EXT selects each action's declared object and hides the undeclared name.	Three fresh KVM boots passed 18/18 physical/logical states, 24/24 Python probes, concurrent 3.10/3.12 views, switch, rollback, and lower-object controls.
Toolchain environments (supporting) [28]	Two installed environments expose distinct interpreter and package trees; one process-group view switches and rolls back.	NAMEI_EXT selects the existing environment root without changing the logical executable path.	Three fresh KVM boots passed 12/12 source states, 12/12 NAMEI_EXT states, 6/6 direct controls, 24/24 stable-root descriptor checks, 12/12 old-descriptor checks, and 36/36 lower-object checks.
Kubernetes ConfigMap publication (supporting) [19, 20]	Official <code>AtomicWriter</code> publication exposes V0, V1, repeated V1, and rollback V0 payloads under one stable root.	NAMEI_EXT selects or hides already-materialized leaf objects while VFS retains modes, ownership, open descriptors, and lower objects.	Three fresh KVM boots completed all nine DMTCP conditions. PathTranslator and NAMEI_EXT each passed three same-path A-to-B transitions; three withdrawn controls returned the expected ENOENT. Restart-time SELECT increased 12-to-24 with the rule and remained 12-to-12 after withdrawal; all 108 before/after lower-object rows were unchanged.
Checkpoint/restore (supporting) [2]	DMTCP restarts one application after its remembered pathname moves from generation A to generation B; withdrawing the mapping must make the reopened path absent.	NAMEI_EXT selects the existing generation directory after restart, while DMTCP retains checkpointing, descriptor restoration, and coordination.	Three fresh KVM boots passed all source, NAMEI_EXT, and withdrawn conditions. All 141 mapping, selection, identity, and preservation rows passed; 204 SELECT events were attributed to the 141 targets, and all three withdrawn-target controls failed with the expected loader diagnostic and no additional target hit.
HPC file staging (supporting) [14]	Pinned LLNL Spindle in serial pull mode runs its unchanged 47-object dynamic-loader slice and populates node-local cache files from the source objects.	NAMEI_EXT selects the exact Spindle-created cache objects for the same loader; Spindle retains cache population and object discovery.	Three fresh KVM boots passed all source, NAMEI_EXT, and withdrawn conditions. All 141 mapping, selection, identity, and preservation rows passed; 204 SELECT events were attributed to the 141 targets, and all three withdrawn-target controls failed with the expected loader diagnostic and no additional target hit.

Table 3: The seven-workload RQ1 answer table. A row becomes evidence only after the source-oracle run passes through the real KVM attach path and lower-filesystem semantics are preserved.

The required RQ1 evidence is a reviewed KVM pass/fail matrix that preserves lookup/readdir coherence and lower-filesystem behavior for each admitted source oracle. A workload is outside NAMEI_EXT's boundary if the oracle needs synthetic contents, data-path mediation, write-conflict handling, or custom metadata ownership in the common path.

The VFS construction invariants are independent of any one source workflow. Table 4 therefore tests the RCU path, registered target lifetime, parent prefilter, and return to ordinary VFS completion directly in the modified-kernel KVM.

Construction invariant	Modified-kernel KVM check	Result
RCU/ref-walk selection	Force <code>RESOLVE_CACHED</code> for a final selected directory and an intermediate selected component.	Both entered selected-target handling from RCU-walk and completed after in-place legitimization, without a full path-walk restart.
Registered-target lifetime	Replace one target ID 128 times while another process continuously reads through the selected pathname.	Every read observed one complete old or new payload.
Parent prefilter	Attach a policy to one exact parent; probe lookup and readdir under that parent and a second, unrelated parent.	The managed parent was filtered and the unrelated parent retained the lower view; a separate zero-invocation control confirmed policy dispatch was bypassed.
Selected-object semantic preservation	Pair direct lower paths with selected paths for a frozen 16-case, 80-operation matrix on ext4/tmpfs. Attribute every selected operation to target selection, PASS, or no policy; retain an opened selected directory across policy teardown.	Three fresh KVM boots passed 48/48 direct and 48/48 selected cases, 240/240 per-operation engagements, six raw identity checks, and 96/96 residual checks. Descriptor-relative create/write/read/rename/unlink completed after teardown with zero subsequent policy engagement.

Table 4: Mechanism checks underlying RQ1. These checks validate construction invariants; they are not additional workload or performance claims.

6.4 RQ2: Cost Versus FUSE

RQ2 isolates the placement cost of policy by comparing `NAMEI_EXT` with a feature-equivalent FUSE policy service under the same oracle. Non-name-resolution lifecycle behavior is delegated identically in both rows. RQ2 does not compare different runtimes. Instead, it compares placing the same lookup/readdir decision inside VFS name resolution with placing it in a FUSE request path.

Workload	Feature-equivalent FUSE comparison	Metrics	Result evidence
Agent workspace lifecycle	Same 48-oracle lifecycle in NAME_EXT and a feature-equivalent FUSE policy service.	Ten paired blocks, 20 KVM boots, 10,000 samples per mechanism.	960/960 oracles passed. Lifecycle p50 was 5.51 μ s versus 62.64 μ s; paired FUSE/NAME_EXT ratio was 11.32 $\times$ [11.24, 11.64].
FxMark cache-hot path walks	Stock, patched-unattached, attached PASS, same-filesystem SELECT, and multithreaded cached FUSE.	50 fresh KVM boots and 450 condition-run observations over nine MRPL/MRPM/MRPH and worker-count cells.	450/450 observations passed. SELECT/FUSE was 1.052–1.088 $\times$, with every paired 95% CI above one; complete SELECT retained 0.895–0.931 of unattached throughput.
Corrected FxMark directory enumeration	Same five conditions over private-directory MRDL and shared-directory MRDM.	Ten paired blocks, 50 fresh KVM boots, and 300 condition-run observations over six operation and worker-count cells.	300/300 observations passed. SELECT/FUSE was 2.20–3.66 $\times$ with CIs above one in five cells; MRDM/4 was 1.018 [0.907, 1.135].
Unused fast path	Stock versus patched-unattached cache-hot MRPL in 30 paired blocks.	60 fresh KVM boots and 180 condition-run observations over three worker-count cells.	Unattached/stock was 1.0009 [0.9921, 1.0036], 1.0083 [0.9950, 1.0179], and 1.0007 [0.9918, 1.0139] at one/two/four workers.

Table 5: RQ2 evidence table. FUSE numbers are interpreted only after the FUSE row implements the same source policy and passes the same source oracle.




Figure 2: Corrected FxMark directory-enumeration throughput. Points are medians over ten fresh-boot blocks; error bars are 95% bootstrap intervals. MRDL gives each worker a private directory; MRDM makes all workers enumerate one shared directory.

The Agent result is a mechanism-level lifecycle comparison, not an end-to-end agent-task speedup. The FxMark FUSE view enabled metadata and kernel caching and issued only 0–19 measured requests per cell; its result does not depend on one daemon round trip per lookup. Attached PASS retained 0.901–0.934 of unattached throughput, while SELECT/PASS was 0.981–0.999. The separate unused-path run supports only cache-hot MRPL on this host; cold caches, other operations, tail latency, and other machines remain open.

Figure 2 broadens RQ2 to the event where the BPF policy runs for every returned directory entry. For private directories, SELECT/FUSE is 2.200–3.663 $\times$, and SELECT scales 3.844 $\times$ from one to four workers versus 2.411 $\times$ for FUSE. Shared-directory enumeration is 2.909 $\times$ and 2.450 $\times$ at one and two workers. At four workers, stock, SELECT, and FUSE converge near five million entries/s; the SELECT/FUSE interval crosses one. The frozen overall verdict is therefore mixed: five cells support the throughput advantage, while shared-directory contention masks it in the sixth.

Evidence	NAMEI_EXT	Wrapfs-derived	Runtime observation	Result
Agent workspace, same 37-row oracle	BPF decides lookup/readdir; ordinary operations use the selected ext4 object.	A port of official Wrapfs registers 34 unique VFS slots across six compiled sources.	Thirteen Wrapfs method classes executed. After NAMEI_EXT detach, an open descriptor completed ordinary lower-file operations without re-entering BPF.	37/37 oracles passed for both mechanisms in each of three KVM boots.
Invalid policy/output matrix	Verifier rejects invalid programs; kernel returns declared errors for malformed or unsupported decisions.	Validation and failures remain kernel-module filesystem responsibilities.	Two verifier cases produced exact rejection logs; 19 independent runtime faults each preserved exact lower-object evidence.	21/21 cells passed in each of three boots.

Table 6: RQ3 boundary table. The comparison is filesystem-method and runtime responsibility for the same oracle-relevant behavior, not a claim that custom filesystems are less expressive.

6.5 RQ3: Ownership And Responsibility Boundary

RQ3 compares ownership boundaries for the same oracle-relevant behavior. The question is not whether custom or stackable filesystems are expressive, because they are. The question is whether the workload needs that broader boundary when the policy effect is only name-resolution selection or visibility.

For this existing-object workspace view, RQ3 is supported: NAMEI_EXT confines workload policy execution to lookup and directory iteration, while the matched stackable implementation owns and executes superblock, lookup, directory, inode, and file methods. This is a workload-specific ownership and responsibility result, not a claim that custom filesystems are unnecessary when an oracle requires synthetic contents, copy-up, conflict resolution, distributed metadata, or data-path mediation.

6.6 Evaluation Scope

The completed evidence includes all seven required formal RQ1 workflows, an Agent lifecycle RQ2 comparison, cache-hot FxMark active and unused-path results, a formal corrected-FxMark directory-enumeration matrix, and one matched Agent-workspace RQ3 boundary experiment. A historical ccache compile matrix is supporting macro evidence: its hot-cache run observed FUSE/NAMEI_EXT 2.18 $\times$ and native/NAMEI_EXT 0.945 $\times$, but it lacks independent-run uncertainty and release-scale miss/stale/corrupt cells. It is not a headline workload because ccache already implements cache lookup and validation in userspace.

The W6 Spindle result covers final pathname selection among 47 cache objects created by the source system; it does not cover Spindle’s distribution, cache-population scalability, or production launch performance. The W4 Kubernetes result covers already-materialized payload selection and visibility, not ConfigMap retrieval, symlink/inotify behavior, or application reload. Cache-cold lookup, broader metadata operations, tail latency, and a second source-derived RQ3 row remain open. Materialized namespaces, eBPF LSM, and native source mechanisms remain related work unless a fixed source oracle makes one a direct comparison.

7 Discussion

7.1 Implications

The broader lesson is that narrow programmable hooks help when applications vary policy but the kernel should retain mechanism ownership. NAMEI_EXT applies that lesson to VFS name resolution rather than to a new filesystem. For NAMEI_EXT, the retained mechanism is VFS path walking and lower-filesystem object semantics. The NAMEI_EXT boundary is useful only when the correctness condition is pathname-to-object selection or directory visibility over existing objects.

FUSE, stackable filesystems, and custom filesystems remain the appropriate implementation boundary when a workload needs filesystem-service features such as synthetic contents, persistent custom metadata, data-path mediation, write-conflict resolution, distributed metadata, or cross-path transactions. NAMEI_EXT is intended to avoid that boundary only for the subset of policies whose effect is visible during lookup or directory enumeration.

7.2 Extension Path

The same boundary also defines future extensions. Synthetic aliases require additional kernel support, while synthetic contents, custom metadata persistence, distributed indexing, and cross-path transactions remain filesystem or runtime responsibilities. `NAMEI_EXT` is intended only for workloads whose correctness oracle is path selection or visibility over existing lower objects.

8 Related Work

Related work falls along a spectrum of filesystem ownership. Namespace mechanisms construct views before lookup, and access-control hooks mediate or observe operations near the VFS. `NAMEI_EXT` puts bounded object-selection policy in VFS name resolution, while FUSE, stackable filesystems, custom filesystems, and metadata services require broader filesystem-service or object-model ownership.

8.1 FUSE Policy Services

FUSE [38], FUSE passthrough [39], ExtFUSE [4], and FUSE-BPF [32] demonstrate the value of placing filesystem policy in, or accelerating, a filesystem-service boundary. ExtFUSE installs thin eBPF request handlers in a FUSE interposition driver and may fall back to its user-space filesystem; the FUSE layer still defines the mounted filesystem’s request handling and cache interaction. FUSE-BPF associates backing objects through a FUSE lookup response and lets BPF filter or alter operations within that stacked filesystem. It can also run passthrough with `root_dir` and `no_daemon`, so daemon elimination is not the architectural distinction. FUSE-BPF still creates a mounted FUSE filesystem, FUSE inodes, and a filesystem-operation forwarding surface that includes lookup, readdir, open, I/O, create, delete, rename, and xattr.

`NAMEI_EXT` instead places the decision in ordinary VFS name resolution. It creates no filesystem instance or FUSE inode: registered targets remain lower-VFS paths, target replacement follows VFS/RCU lifetime rules, and successful selection exits the extension and resumes ordinary permission and open completion. Its static-key and parent filter also define cost for processes and parents outside the policy scope.

RFUSE, CoFS, and DFUSE further show that the FUSE boundary can be optimized or specialized with scalable kernel-userspace communication, fixed container-image lookup structures, and distributed-cache consistency support [9, 47, 21]. These systems make FUSE the RQ2 baseline family, not a strawman: a fair comparison must implement the same path-view policy under the same source oracle and justify the FUSE caching/passthrough configuration. `NAMEI_EXT` applies only when the source oracle needs selection or visibility over existing objects and does not need the broader filesystem service.

8.2 Custom, Stackable, And Metadata-Service Filesystems

Wrapfs [52], Bento [25], and custom filesystems demonstrate broader filesystem-extension boundaries. They can implement richer behavior and, in Bento’s case, reduce the risk of in-kernel filesystem development. They are the RQ3 ownership comparison point because the developer still owns filesystem methods, storage semantics, or framework-specific filesystem behavior. These mechanisms remain the appropriate implementation boundary when the workload needs custom file operations, synthetic contents, persistence, or richer semantics.

DeltaFS, IndexFS, and TableFS show that full metadata services and specialized filesystems are valuable when the problem is metadata distribution, indexing, or storage layout [54, 31, 30]. Recent systems reinforce the same boundary: Oxbow coordinates kernel, userspace, and device components; DeLFS builds a decentralized kernel filesystem; FalconFS moves path resolution and metadata indexing into a distributed filesystem service; SwitchFS and MesaFS target distributed metadata updates; and SYSSPEC/SpecFS generates a complete filesystem from specifications [17, 1, 50, 49, 15, 24]. They are not the main workload baselines for `NAMEI_EXT`, but they clarify the boundary because such workloads require a metadata service or full filesystem when the oracle needs distributed metadata, data-path ownership, persistence, or a new filesystem object model.

8.3 Namespace Construction

Bind mounts, projected volumes, symlink forests, copy trees, and OverlayFS are practical ways to present alternate views [19, 43]. They preserve ordinary kernel filesystem behavior, but express policy by constructing or updating namespace state. `NAMEI_EXT` asks when the policy can stay at lookup and directory-enumeration time instead of materializing a view before the transition completes.

8.4 Access Control And Observation Hooks

Landlock, BPF LSM programs, and fanotify show that Linux already exposes specialized hooks around filesystem access and events [41, 42, 22]. Among neighboring mechanisms, these hooks sit between materialized namespace construction and `NAMEI_EXT`. They can run verified policy in the kernel, but their natural role is to deny, mediate, or observe operations. `NAMEI_EXT` selects which existing object a name denotes during resolution and then returns control to the ordinary VFS and lower filesystem. Access-control denial remains outside the current action set because the primary claim is name-resolution selection and visibility. USEC is a recent access-control neighbor with an LSM-compatible mandatory-access-control design [16]; XRP, vBPF, PeeR, KRAKENGUARD, Xkernel, and bpftime/EIM are adjacent programmability work on storage-function hooks, eBPF virtualization, eBPF scheduling, eBPF isolation, kernel tunability, and userspace extension safety [56, 53, 5, 29, 8, 55]. They help set safety and overhead expectations for programmable kernel extensions, but they do not provide a VFS name-resolution object-selection boundary.

8.5 Source Workload Systems

Agent filesystems are closest to our workloads, but they are workload sources rather than direct cost baselines [46, 26, 27, 51]. They show that agent workspaces need branch, sandbox, staging, and snapshot-like views, while their implementations also own behavior beyond `NAMEI_EXT`'s boundary. For agent workloads, custom or stackable filesystem ownership provides the RQ3 responsibility comparison, while FUSE remains the RQ2 cost comparison. `NAMEI_EXT` applies when the same source oracle can be satisfied by changing path visibility or object selection while leaving the rest of that ownership with the source system or lower filesystem.

Environment and cache systems provide a second workload source, not a filesystem mechanism baseline. SWE-Factory-Gym also supplies the released source-task oracle used in the Agent workspace case. Along with MEnvData-SWE and SWE-rebench V2, it provides build/test oracles for environment state and cache reuse policies [11, 33, 13, 34]. CoFS and large-scale LLM data-pipeline work add contemporary evidence that build, container, and data-cache paths matter to real systems [47, 7]. Agentic systems such as Murakkab and SREGym show that modern agents run through multi-step tool, workflow, and live-environment oracles [6, 10]. These sources motivate workload state and oracle selection. Their runtime orchestration, environment synthesis, projected-volume materialization, and reload behavior remain outside `NAMEI_EXT`'s mechanism claim.

9 Conclusion

`NAMEI_EXT` frames state-dependent path views as a systems problem in which pathname policy is programmable while filesystem ownership remains with the VFS and lower filesystem. The abstraction occupies the space between namespace materialization and full programmable filesystems, following the same policy-versus-ownership split used by `sched_ext`. The design and prototype show that lookup-time policy can be exposed as one bounded eBPF decision while VFS and lower filesystems retain object ownership. The evaluation therefore centers on three lookup/readdir claims: expressiveness for source-derived path-view policies, placement cost versus feature-equivalent FUSE, and method and runtime responsibility versus custom or stackable filesystems.

References

- [1] Taehwan Ahn, Chanhyeon Yu, Sangjin Lee, and Yongseok Son. DeLFS: A decentralized Log-Structured file system for manycores. In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pages 793–807, Seattle, WA, July 2026. USENIX Association.
- [2] Jason Ansel, Kapil Arya, and Gene Cooperman. Transparent checkpointing of the whole software stack: the dmtcp approach. In *Proceedings of the 2016 Workshop on Fault Tolerance for HPC at Extreme Scale*, 2016.
- [3] Bazel. Sandboxing. <https://bazel.build/docs/sandboxing>. Accessed 2026-06-14.
- [4] Ashish Bijlani and Umakishore Ramachandran. Extension framework for file systems in user space. In *USENIX ATC*, 2019.
- [5] Jeremy Carin, Ben Holmes, Weiyang Wang, Ankit Bhardwaj, and Manya Ghobadi. PeeR: First-Class scheduling for Latency-Critical eBPF applications. In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pages 2465–2482, Seattle, WA, July 2026. USENIX Association.
- [6] Gohar Irfan Chaudhry, Esha Choukse, Haoran Qiu, Inigo Goiri, Rodrigo Fonseca, Adam Belay, and Ricardo Bianchini. Murakkab: Resource-Efficient agentic workflow orchestration in cloud platforms. In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pages 567–587, Seattle, WA, July 2026. USENIX Association.
- [7] Luofan Chen, Chenhan Wang, Weidong Zhang, Jinxin Chi, Hequan Zhang, Zhanbo Wang, Chenyuan Wang, Lishu Luo, Sijin Wu, Junqi Hu, Jun Wang, Cheng Chen, Lixin Huang, Liyang Zhao, Yong Tian, Jun Guo, Youhui Bai, Wencong Xiao, Kang Chen, and Cheng Li. Teaching the old dog new tricks: Building efficient data pipelines for Large-Scale LLM Pre-Training (operational systems). In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pages 349–365, Seattle, WA, July 2026. USENIX Association.
- [8] Zhongjie Chen, Wentao Zhang, Yulong Tang, Ran Shu, Fengyuan Ren, Tianyin Xu, and Jing Liu. Xkernel: Principled performance tunability of operating system kernels. In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pages 547–565, Seattle, WA, July 2026. USENIX Association.
- [9] Kyu-Jin Cho, Jaewon Choi, Hyungjoon Kwon, and Jin-Soo Kim. RFUSE: Modernizing userspace filesystem framework through scalable Kernel-Userspace communication. In *22nd USENIX Conference on File and Storage Technologies (FAST 24)*, pages 141–157, Santa Clara, CA, February 2024. USENIX Association.
- [10] Jackson Clark, Yiming Su, Saad Mohammad Rafid Pail, Yifang Tian, Lily Gniedziejko, Hans-Arno Jacobsen, Yinfang Chen, and Tianyin Xu. SREGym: A live benchmark for AI SRE agents with High-Fidelity failure scenarios. arXiv:2605.07161, 2026.
- [11] DeepSoftwareAnalytics. Swe-factory. <https://github.com/DeepSoftwareAnalytics/swe-factory>, 2025. Accessed 2026-07-10.
- [12] Docker Docs. What is docker? <https://docs.docker.com/get-started/docker-overview/>, 2026. Accessed 2026-07-13.
- [13] ERNIE Research. Menvagent and menvdata-swe. <https://github.com/ernie-research/MEnvAgent>, 2026. Accessed 2026-07-10.
- [14] Wolfgang Frings, Dong H. Ahn, Matthew LeGendre, Todd Gamblin, Bronis R. de Supinski, and Felix Wolf. Massively parallel loading. In *Proceedings of the 27th International ACM Conference on International Conference on Supercomputing*, pages 389–398. ACM, 2013.
- [15] Hao Guo, Jiwu Shu, and Youyou Lu. MesaFS: An I/O-Efficient metadata service for distributed file systems. In *Proceedings of the 21st European Conference on Computer Systems (EuroSys 26)*, 2026.
- [16] Yu Jiang, Wenhuan Liu, Fuchen Ma, Yuheng Shen, Yuanliang Chen, Lei Zhang, He Li, Quan Zhang, and Chijin Zhou. USEC: A User-Requirement-Driven mandatory access control framework for operating systems (operational systems). In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pages 191–206, Seattle, WA, July 2026. USENIX Association.
- [17] Jongyul Kim, Jaehwan Lee, Inhoe Koo, Peizhe Liu, Jiyuan Zhang, Junho Ahn, Tianyin Xu, and Youngjin Kwon. Oxbow: A coordinated architecture for Multi-Component file systems. In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pages 1425–1442, Seattle, WA, July 2026. USENIX Association.
- [18] Kubernetes. Secrets. <https://kubernetes.io/docs/concepts/configuration/secret/>. Accessed 2026-06-14.
- [19] Kubernetes. Projected volumes. <https://kubernetes.io/docs/concepts/storage/projected-volumes/>, 2026. Accessed 2026-06-14.

- [20] Kubernetes documentation. Configure a pod to use a configmap. <https://kubernetes.io/docs/tasks/configure-pod-container/configure-pod-configmap/>, 2026. Accessed 2026-07-13.
- [21] Haoyu Li, Jingkai Fu, Qing Li, Windsor Hsu, and Asaf Cidon. DFUSE: Strongly consistent Write-Back kernel caching for distributed userspace file systems. arXiv:2503.18191, 2025.
- [22] Linux man-pages. fanotify(7). <https://man7.org/linux/man-pages/man7/fanotify.7.html>, 2026. Accessed 2026-06-14.
- [23] Linux man-pages. mount(2) linux manual page. <https://man7.org/linux/man-pages/man2/mount.2.html>, 2026. Accessed 2026-07-13.
- [24] Qingyuan Liu, Mo Zou, Hengbin Zhang, Dong Du, Yubin Xia, and Haibo Chen. Sharpen the spec, cut the code: A case for generative file system with SYSSPEC. In *24th USENIX Conference on File and Storage Technologies (FAST 26)*, pages 291–311, Santa Clara, CA, February 2026. USENIX Association.
- [25] Samantha Miller, Kaiyuan Zhang, Mengqi Chen, Ryan Jennings, Ang Chen, Danyang Zhuo, and Thomas Anderson. High velocity kernel file systems with bento. In *FAST*, 2021.
- [26] MultiKernels. Branchfs. <https://github.com/multikernel/branchfs>, 2026. Accessed 2026-07-10.
- [27] MultiKernels. Sandlock. <https://github.com/multikernel/sandlock>, 2026. Accessed 2026-07-10.
- [28] Nix. Nix store. <https://nix.dev/manual/nix/2.34/store/>. Accessed 2026-06-14.
- [29] Jainil Patel, Lucas Graeff Buhl-Nielsen, Adrien Ghosn, and Marios Kogias. KRAKENGUARD: Towards Fine-Grained eBPF isolation. In *23rd USENIX Symposium on Networked Systems Design and Implementation (NSDI 26)*, pages 2685–2704, Renton, WA, May 2026. USENIX Association.
- [30] Kai Ren and Garth Gibson. Tablefs: Enhancing metadata efficiency in the local file system. In *USENIX ATC*, 2013.
- [31] Kai Ren, Qing Zheng, Swapnil Patil, and Garth Gibson. Indexfs: Scaling file system metadata performance with stateless caching and bulk insertion. In *SC '14: Proceedings of the International Conference for High Performance Computing, Networking, Storage and Analysis*, 2014.
- [32] Daniel Rosen. FUSE BPF: A stacked filesystem extension for fuse. <https://patchew.org/linux/20230418014037.2412394-1-drosen%40google.com/>, 2023. Accessed 2026-07-14.
- [33] SWE-Factory. SWE-Factory-Gym dataset. <https://huggingface.co/datasets/SWE-Factory/SWE-Factory-Gym>, 2026. Dataset row `pallets_click-2622`; accessed 2026-07-29.
- [34] SWE-rebench. Swe-rebench v2. <https://github.com/SWE-rebench/SWE-rebench-V2>, 2026. Accessed 2026-07-10.
- [35] The Linux Kernel documentation. Bpf documentation. <https://docs.kernel.org/bpf/>, 2026. Accessed 2026-07-13.
- [36] The Linux Kernel documentation. ebpf verifier. <https://docs.kernel.org/bpf/verifier.html>, 2026. Accessed 2026-07-13.
- [37] The Linux Kernel documentation. Extensible scheduler class. <https://docs.kernel.org/scheduler/sched-ext.html>, 2026. Accessed 2026-06-14.
- [38] The Linux Kernel documentation. Fuse overview. <https://docs.kernel.org/filesystems/fuse/fuse.html>, 2026. Accessed 2026-06-14.
- [39] The Linux Kernel documentation. Fuse passthrough. <https://docs.kernel.org/next/filesystems/fuse-passthrough.html>, 2026. Accessed 2026-06-14.
- [40] The Linux Kernel documentation. Kvm. <https://docs.kernel.org/virt/kvm/index.html>, 2026. Accessed 2026-07-13.
- [41] The Linux Kernel documentation. Landlock: unprivileged access control. <https://docs.kernel.org/userspace-api/landlock.html>, 2026. Accessed 2026-06-14.
- [42] The Linux Kernel documentation. Lsm bpf programs. https://docs.kernel.org/bpf/prog_lsm.html, 2026. Accessed 2026-06-14.
- [43] The Linux Kernel documentation. Overlay filesystem. <https://docs.kernel.org/filesystems/overlayfs.html>, 2026. Accessed 2026-06-14.
- [44] The Linux Kernel documentation. Overview of the linux virtual file system. <https://docs.kernel.org/filesystems/vfs.html>, 2026. Accessed 2026-07-13.

- [45] The Linux Kernel documentation. Pathname lookup. <https://docs.kernel.org/filesystems/path-lookup.html>, 2026. Accessed 2026-07-13.
- [46] Turso Database. Agentfs. <https://github.com/tursodatabase/agentfs>, 2026. Accessed 2026-07-10.
- [47] Li Wang, Jinxu Du, Yang Yang, Qingbo Wu, Tao Liu, and Haoze Wu. CoFS: A filesystem for fast container startup. In *24th USENIX Conference on File and Storage Technologies (FAST 26)*, pages 415–423, Santa Clara, CA, February 2026. USENIX Association.
- [48] XDG Desktop Portal Project. Documents portal. Official API documentation, 2026.
- [49] Jingwei Xu, Mingkai Dong, Qiulin Tian, Ziyi Tian, Tong Xin, and Haibo Chen. SwitchFS: Asynchronous metadata updates for distributed filesystems with In-Network coordination. EuroSys 2026; arXiv:2410.08618, 2026.
- [50] Jingwei Xu, Junbin Kang, Mingkai Dong, Mingyu Liu, Lu Zhang, Shaohong Guo, Ziyang Qiu, Mingzhen You, Ziyi Tian, Anqi Yu, Tianhong Ding, Xinwei Hu, and Haibo Chen. FalconFS: Distributed file system for Large-Scale deep learning pipeline. In *23rd USENIX Symposium on Networked Systems Design and Implementation (NSDI 26)*, pages 431–447, Renton, WA, May 2026. USENIX Association.
- [51] YoloFS. YoloFS. <https://github.com/YoloFS/YoloFS>, 2026. Accessed 2026-07-10.
- [52] Erez Zadok, Ion Badulescu, and Alex Shender. Extending file systems using stackable templates. In *USENIX ATC, FREENIX Track*, 1999.
- [53] Jing Zhang, Xiaguannan Song, Dong Du, Yubin Xia, Binyu Zang, and Haibo Chen. Virtualizing eBPF with Late-Binding. In *20th USENIX Symposium on Operating Systems Design and Implementation (OSDI 26)*, pages 2127–2148, Seattle, WA, July 2026. USENIX Association.
- [54] Qing Zheng, Charles D. Cranor, Gregory R. Ganger, Garth A. Gibson, George Amvrosiadis, Bradley W. Settlemyer, and Gary A. Grider. Deltafs: A scalable no-ground-truth filesystem for massively-parallel computing. In *SC*, 2021.
- [55] Yusheng Zheng, Tong Yu, Yiwei Yang, Yanpeng Hu, Xiaozheng Lai, Dan Williams, and Andi Quinn. Extending applications safely and efficiently. In *19th USENIX Symposium on Operating Systems Design and Implementation (OSDI 25)*, pages 557–574, Boston, MA, July 2025. USENIX Association.
- [56] Yuhong Zhong, Haoyu Li, Yu Jian Wu, Ioannis Zarkadas, Jeffrey Tao, Evan Mesterhazy, Michael Makris, Junfeng Yang, Amy Tai, Ryan Stutsman, and Asaf Cidon. XRP: In-Kernel storage functions with eBPF. In *16th USENIX Symposium on Operating Systems Design and Implementation (OSDI 22)*, pages 375–393, Carlsbad, CA, July 2022. USENIX Association.